

Problem Statement:

An executive assistant (Flow), which will prepare you before every meeting.

Busy professionals often rush between meetings without time to gather key information. Important details—like past emails, CRM notes, shared files, and web insights—are scattered across different systems. This fragmentation forces executives, sales teams, and project leads to scramble for information, leading to unprepared discussions, misaligned expectations, and lost productivity. By automatically compiling and distilling relevant data into a concise pre-meeting briefing, we eliminate tedious manual prep, reduce context-switching overhead, and empower more focused, strategic conversations. Solving this problem is vital: executives report feeling under-prepared for nearly 40% of their 23 weekly meeting hours, costing organizations significant time, money, and decision quality.

Target Audience & Context:

Our targeted users are high-value professionals like- **C-Suite executives, account managers, consultants, and project leads**- who have lots of meetings. They need up-to-date information (like recent interactions, open tasks, and customer feedback) to work effectively. In remote or hybrid work, they miss out on quick, informal updates, so automatic briefings are even more important. Our solution connects with the tools they already use (like Google/ Outlook Calendar, Gmail/ Outlook, Salesforce/ HubSpot, Slack/ Teams) to provide timely insights with minimal effort, making it easy to adopt and use.

Use of Gen-AI:

Generative AI powers the core intelligence of our briefing assistant. A fine-tuned large language model (e.g., ChatGPT or Gemini AI) ingests raw inputs which are email bodies, CRM records, document snippets, and relevant web data and then produces a human-readable briefing.

Key AI capabilities include:

Contextual Summarization: Condensing lengthy email threads into bullet-point recaps that highlight decisions, deadlines, and sentiment.

Entity Recognition & Linking: Identifying people, companies, and projects, then correlating them across disparate sources for unified context.

Dynamic Enrichment: Querying news APIs or knowledge graphs to append up-to-date industry or competitor insights.

Template-Driven Output: Using prompt engineering to ensure consistent formatting—sections for summary, last touchpoint, and open action items—facilitating rapid comprehension.

Solution Framework:

Data Ingestion & Authentication

- **Calendar API:** Retrieve upcoming meeting metadata (title, time, participants).
- **Email Connectors:** Secure OAuth2 access to fetch recent threads (Gmail/Outlook).
- **CRM API:** Query Salesforce/HubSpot for contact histories, deal stages, and notes.
- **Document Store:** Access Drive/SharePoint for relevant decks and proposals.
- **Web Data:** Pull public updates via NewsAPI or lightweight scrapers.

Preprocessing & Correlation

- **Normalization:** Standardize names, dates, and identifiers.
- **Deduplication:** Remove repeated content across sources.
- **Relevance Filtering:** Prioritize items from the past 30 days and high-impact interactions.

Gen-AI Summarization Engine

- **Prompt Pipeline:** Sequential prompts guide the model to extract agenda, summarize last interaction, and list pending tasks.
- **Entity Linking:** Tag and hyperlink stakeholders to CRM profiles.
- **Sentiment Flags:** Surface any negative or urgent tones for immediate attention.

Brief Assembly & Delivery

- **Template Renderer:** Populate a one-page briefing.
- **Notification Service:** Dispatch via email or post snippet to Slack/Teams 15 minutes before meetings.

Security & Compliance

- **Encryption:** TLS in transit, AES-256 at rest.
- **Access Controls:** Least-privilege OAuth scopes, periodic token rotation.
- **Audit Logging:** Track all data fetches and summary generations.

Feasibility & Execution:

In a 12–24-hour sprint with a 4 members team, we'll build an MVP using low-code connectors (e.g., Zapier or n8n) for API integration, OpenAI's GPT-4 API for summarization, and a lightweight Node.js/Express backend. A minimal React or Slack-bot frontend will handle notifications. Dependencies (calendar/email SDKs, CRM REST APIs, NewsAPI) are off-the-shelf, enabling rapid assembly. By focusing on core flows—calendar fetch, AI-driven summary, and Slack delivery—we'll validate technical feasibility and user value before iterating to add document ingestion and advanced filters.

Scalability & Impact:

Our modular microservices architecture (containerized via Docker/Kubernetes) supports horizontal scaling: ingestion, AI engine, and delivery layers scale independently. New data sources (e.g., project management tools) plug in with minimal development. As adoption grows, caching, rate-limiting, and prioritized queues ensure performance. Recovering just two prep hours per executive per week amounts to hundreds of hours saved per team per quarter. Enhanced meeting readiness accelerates deal closures, reduces follow-ups, and elevates decision quality, driving measurable ROI across the organization.

Conclusion & Bonus Minimum Lovable Product:

Our AI-driven Executive Assistant automates pre-meeting prep, transforming fragmented data into clarity and strategic focus. Within 24 hours, the MVP validates core functionality. To become a lovable product, we'll launch a freemium model: basic calendar-email briefs for free, with premium multi-source integrations, analytics dashboards, and a connector marketplace—unlocking rapid user growth, viral expansion, and monetization potential.

Team Flow

